

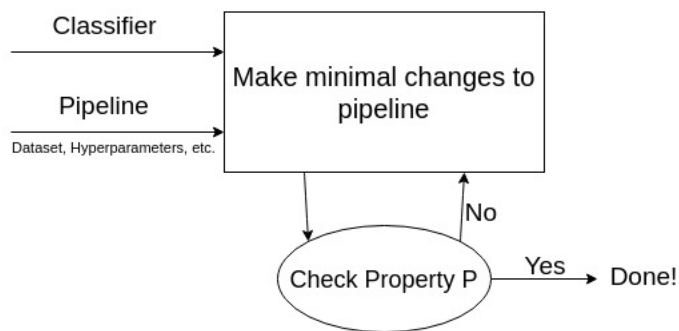
CS-799: Notes

Goutham Ramakrishnan

Guide: Prof. Aws Albarghouthi

Broad goal:

Given a machine learning pipeline with a dataset and classifier, find a minimal set of changes which need to be made to the pipeline for it to satisfy a given property P.



Ideas:

- Use clustering (K-Means? Hierarchical? Semi-supervised?)
- Make changes to dataset iteratively
- Changes made must be interpretable
- Use training data debugging using trusted items as Property P to be starting point for exploring approaches

Discussion with Xuezhou:

- Using clustering as starting point for DUTI
- Active Debugging: Prof. Po-Ling's student working on it. Requires periodic expert feedback.
- For interpretability, find bugs first and then learn **rule list**
- Debugging the machine learning pipeline: Talk by Prof. Jerry Zhu (<http://pages.cs.wisc.edu/jerryzhu/adversarial/pub/debugML.pdf>)
- Difficult to test approach due to lack of good datasets for this kind of task

Thoughts on clustering:

- Advantage of hierarchical clustering over K-means: The output hierarchy is more informative than output of K-means, number of clusters need not be pre-specified, can use dendrogram to choose K. Tradeoff: Speed and efficiency
- Review on Semi-supervised clustering methods: <https://arxiv.org/pdf/1307.0252.pdf>
- Constrained clustering looks promising, can enforce cannot-link constraints on differently labeled trusted items
- Need to read up more on semi-supervised hierarchical clustering

Formulation:

Training set: $(X, Y) = \{(x_i, y_i)\}_{1:n}$

Trusted set: $(X, Y) = \{(\tilde{x}_i, \tilde{y}_i)\}_{1:m}$

Number of clusters: K, Set of clusters $S = \{S_1, \dots, S_K\}$

K-Means:

$$\operatorname{argmax}_S \sum_{k=1}^K \sum_{x \in S_i} \|x - \mu_k\|_2^2 = \operatorname{argmin}_S \sum_{k=1}^K |S_k| \operatorname{Var}(S_i)$$

To minimize variance, add:

$$\sum_{k=1}^K \operatorname{Entropy}(y_i | x_i \in C_k)$$

To ensure trusted items with different labels are not in the same cluster, add:

$$\sum_{\tilde{x}_i, \tilde{x}_j: \tilde{y}_i \neq \tilde{y}_j} \operatorname{Cost}(\tilde{x}_i, \tilde{x}_j) \mathbb{1}(C(\tilde{x}_i) = C(\tilde{x}_j))$$

$C(\tilde{x}_i)$ is the cluster to which \tilde{x}_i belongs.

Notes: October 26

- Finding the ideal number of clusters K for a dataset is an unsolved problem, and there is no universally agreed upon metric upon which K can be decided. Methods include the elbow method, silhouette analysis, gap statistic (define a maximum number of clusters choose K with highest gap statistic), CH index etc.
- **Gap statistic:** Can be used for automatically deciding number of clusters. Compares average distance between points in a cluster with respect to a reference dataset (drawn from uniform distribution of features).
(<http://www.sthda.com/english/wiki/print.php?id=239>)
- Hierarchical clustering cannot be used for "predicting" the cluster of a new point (Implying that in our current procedure, trusted items need to be added to the buggy dataset before clustering). An alternative strategy is to use Agglomerative Clustering, and then train another classifier (like KNN) that can do predictions. However, clustering by itself is not a means of classification, and is hard to justify (as each new data point may potentially modify the clustering).
- **Choosing K using dendrograms:** Dendrograms are a visual heuristic, and are not practical to view when the number of data points is large. Moreover, inspecting dendrograms and deciding K is also a subjective process.
- Clustering as a concept is difficult to justify when the features are categorical. The primary hurdle then becomes defining an appropriate "distance" between points. This can be highly dataset dependent.
- Two possibilities are the generalized Mahalanobis distance and the Gower distance. With these distance measures, KMeans algorithm cannot be used directly (as the "centroid" of each cluster ceases to have any meaning). Hierarchical clustering is our best bet in this regard.
- **Gower distance:** General formula for estimating "distances" between datapoints when features are a combination of categorical and numeric.
(<http://halweb.uc3m.es/esp/Personal/personas/jmmarin/esp/MetQ/Talk6.pdf>)
- **Choices of Linkage in Agglomerative Clustering:** The three simplest linkage strategies are complete, single and average which minimize the maximum, minimum and average "distances" between points in two clusters. All three give very different clusterings in the end. Single linkage in particular suffers from chaining (too many points in one cluster), this also occurred during my experiments with the GermanLoan dataset. (Complete linkage worked acceptably)
- **Experiments:** Conducted experiments on German Loan dataset using Gower distance, agglomerative clustering with complete linkage. Dendrogram is hard to interpret. The distribution of trusted items is across clusters, flipping one cluster at a time did not lead to any conclusive results.